

# Introduction to Project 3

Discussion 7

# Project 3 Goals

- Understanding the ARIES recovery algorithm
- Implementing the ARIES recovery algorithm in C++

# Getting Started

- Read the recovery sections of the textbook:
  - Section 16.7 – Introduces the recovery manager
  - Chapter 18 – Covers the ARIES recovery Algorithm
    - Super important to read and understand this section because this is what the project wants you to implement

# Getting Started

- Download the zip file contained in the project description on Canvas
- Understand the different components of the recovery simulator
  - LogRecord
  - Storage Engine
  - Main
  - LogMgr
    - This is what you will implement for the project

# Grading

- You can access the autograder at <https://grader484.eecs.umich.edu/>
- Autograder beta version is now open
- Limit of 4 submission per day
- Autograder runs hidden test suite using your LogMgr.cpp and checks the output
- Highest score will be the final grade
  - If you submit late, the penalty is only applied to the late submission

# The Disk

- The disk is represented as a simple text file
- Each line corresponds to a page
- Each page consists of an LSN followed by a string which is the data stored on the page
- An example of the disk is located in `StorageEngine/sampleDBFile.txt`

# Testcases

- testcases/ contains 6 basic test cases to help check that your project is working
- Passing all 6 tests cases does not mean you will pass the autograder. Create your own test cases as well.
- Results of testcases can be found in output/log and output/dbs/
  - Remember to remove the files if you run the case again

# Testcase Format

- First line of the file tells the simulator where to find the text file representing the initial state of the disk
- All lines after specify an action to perform
  - Write - txid, action, pageid, offset, data
  - Commit
  - Abort – txid, action, # of writes allowed during abort
  - Checkpoint – action
  - Crash – action {# of writes allowed during crash}
  - end marks the end of the transaction



# Running Project 3

- In the parent direction, run make
  - You may want to modify the Makefile to do more complex operations such as running all the testscases
- From the parent direction, run the simulator with the command: `./main.o testcases/test00`
  - This will run test00

# LogRecord.h

- Contains txTableEntry struct
  - Create new entries by calling txTableEntry(lsn, status) constructor
- Contains four log record classes:
  - LogRecord
    - Base log record type
  - UpdateLogRecord
    - Log record for updates
  - CompensationLogRecord
    - Log Record for CLR
  - ChkptLogRecord
    - Log Record for checkpoint operations

# LogRecord.cpp

- Two functions we care about here:
  - toString() – This will convert a log record of any type into a string.
  - stringToRecordPtr – This will convert a string corresponding to a single log record into a LogRecord pointer

# Storage Engine

- This keeps track of pages and manages disk access
  - Adds pages to in-memory buffer if read/write performed
  - Flushes pages to memory if buffer is full
- A page consists of an id, an LSN, a dirty bit, and a string of data.

# Storage Engine

- Functions that may be useful:
  - nextLSN() - This will return the a unique ID assigned in monotonically increasing order
  - pageWrite(...) – This will write a page into memory
    - This can return false if you are no longer allowed to write pages. This is determined by `page_write_permitted`.
    - Normally the storage engine will do this, but this is needed for aborts or recovery

# Storage Engine

- Functions that may be useful:
  - getLSN(...) - This will return the lastLSN of a the given page id
  - store\_master(...) – This writes an int to a location in disk
  - get\_master() – This retrieves the int written by store\_master
  - getLog() – Gets the log from disk

# Main

- This reads in the testcase files and initializes the log manager and storage engine
- It also parses the testcase file and calls the appropriate functions

# LogMgr

- This is the file you are turning in for the project
- Implement all the functions specified in LogMgr.h
- You must maintain the transaction table and the dirty page table. These are stored as maps within LogMgr.h



# Debugging the Code

- Using gdb on the CAEN machine will probably have issues displaying certain data structures like maps.
- One possible solution :
  - mkdir <directory\_name>
  - cd <directory\_name>
  - svn co svn://gcc.gnu.org/svn/gcc/branches/gcc-4\_6-branch/libstdc++-v3/python
  - vim ~/.gdbinit
  - Add the code on the next slide to gdbinit

# Code to add to gdbinit

- Replace /path/to/<directory\_name>/python with the actual filepath
- python  
import sys  
sys.path.insert(0, '/path/to/gdb\_printers/python')  
from libstdcxx.v6.printers import  
register\_libstdcxx\_printers  
register\_libstdcxx\_printers (None)  
end

# Things to note

- **DO NOT** modify any file other than LogMgr.cpp
- **DO NOT** directly read/write to disk. You must use the StoreEngine interface to do this
- **DO NOT** create static variables or functions in LogMgr.cpp
- **DO** get started early. Get something ready for when the autograder opens
- **DO** read the textbook chapters. Chapter 18 contains most of the information you need to do the project

# Things to note

- In the stringToLRVector function, add:

```
vector<LogRecord*> result;
```

```
istringstream stream(logstring);
```

```
string line;
```

```
while(getline(stream, line)) {
```

```
    LogRecord* lr =
```

```
    LogRecord::stringToRecordPtr(line);
```

```
    Results.push_back(lr);
```

```
}
```

```
return result;
```

# Questions about the Project?

# Example

- 1 write 2 0 one
- 2 write 3 0 two
- 2 commit
- 3 write 1 0 three
- 3 write 2 0 four
- crash {7}

# Stable Storage

Disk State

-1 XXXXXXXXXXXX

-1 XXXXXXXXXXXX

-1 XXXXXXXXXXXX

-1 XXXXXXXXXXXX

Log State

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After

# T1 writes

- 1 write 2 0 one

Transaction Table

transl D	lastLS N	Stat
1	2	U

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one

Dirty Page Table

pageID	recLSN
2	2



# T2 writes

- 2 write 3 0 two

Transaction Table

transl D	lastLS N	Stat
1	2	U
2	3	U

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two

Dirty Page Table

pageID	recLSN
2	2
3	3

# T2 commits

- 2 commit

Transaction Table

transl D	lastLS N	Stat
1	2	U
2	3	U

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				

Dirty Page Table

pageID	recLSN
2	2
3	3

# Stable Storage – After commit

Disk State

-1 XXXXXXXXXXXX

-1 XXXXXXXXXXXX

-1 XXXXXXXXXXXX

-1 XXXXXXXXXXXX

Log State							
LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				

# After Commit

- 2 commit

Transaction Table

transl D	lastLS N	Stat
1	2	U

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
5	4	2	end				

Dirty Page Table

pageID	recLSN
2	2
3	3

# T3 Writes

- 3 write 1 0 three

Transaction Table

transl D	lastLS N	Stat
1	2	U
3	6	U

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
5	4	2	end				
6	-1	3	update	1	0	xxxxx	three

Dirty Page Table

pageID	recLSN
2	2
3	3
1	6

# T3 Writes

- 3 write 2 0 four

Transaction Table

transl D	lastLS N	Stat
1	2	U
3	7	U

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
5	4	2	end				
6	-1	3	update	1	0	xxxxx	three
7	6	3	update	2	0	onex	four

Dirty Page Table

pageID	recLSN
2	2
3	3
1	6

# Crash

- crash {7}

Transaction Table

transl D	lastLS N	Stat

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After

Dirty Page Table

pageID	recLSN

# Analysis Steps

- Find the most recent begin checkpoint. If there isn't one, then start from the beginning of the log
- Look at each record:
  - End log: Remove T from the transaction table
  - Others: Add T to the transaction table if not there.

Change lastLSN

If it is a commit, set status to C.

If it is a redoable record on P, add P to the dirty page table



# Analysis

- crash {7}

Transaction Table

transl D	lastLS N	Stat

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After

Dirty Page Table

pageID	recLSN

Disk Log

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				

# Analysis

- crash {7}

Transaction Table

transl D	lastLS N	Stat
1	2	U

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After

Dirty Page Table

pageID	recLSN
2	2

Disk Log

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				

# Analysis

- crash {7}

Transaction Table

transl D	lastLS N	Stat
1	2	U
2	3	U

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After

Dirty Page Table

pageID	recLSN
2	2
3	3

Disk Log

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				

# Analysis

- crash {7}

Transaction Table

transl D	lastLS N	Stat
1	2	U
2	4	C

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After

Dirty Page Table

pageID	recLSN
2	2
3	3

Disk Log

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				

# Redo Steps

- Find oldest update in log (smallest recLSN) and start at that point in the log
- For each redoable record:
  - Is the page in dirty page table?
  - The page is in dirty page table, but is recLSN for the page is less than or equal to the LSN of the record?
  - Is the pageLSN less than the LSN of the log record?
  - If it yes for all three, redo the record
- Remove committed transactions from table

# Redo

- crash {7}

Transaction Table

transl D	lastLS N	Stat
1	2	U
2	4	C

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After

Dirty Page Table

pageID	recLSN
2	2
3	3

Disk Log

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				

# Redo

- crash {7}

Transaction Table

transl D	lastLS N	Stat
1	2	U
2	4	C

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After

Dirty Page Table

pageID	recLSN
2	2
3	3

Disk Log

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				

# Redo

- crash {7}

Transaction Table

transl D	lastLS N	Stat
1	2	U

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
8	4	2	end				

Dirty Page Table

pageID	recLSN
2	2
3	3

Disk Log

LSN	prvLSN	transID	type	pageID	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				



# Stable Storage – After redo

Disk State

-1 XXXXXXXXXXXX

2 XXXXXXXXXXXX

3 XXXXXXXXXXXX

-1 XXXXXXXXXXXX

Log State							
LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				

# Undo Steps

- Undo all in the transaction table starting with the transaction with the largest LSN value in transaction table
- For each record:
  - If CLR:
    - If undoNextLSN is null, add end record to log
    - Otherwise, add undoNextLSN to the set to undo
  - If update:
    - Create a CLR record in the log. Add end record if undoNext is null
    - Add prevLSN to set to undo

# Undo

- crash {7}

Transaction Table

transl D	lastLS N	Stat

Log tail

LSN	prvLSN	transID	type	pageID or undoNext	offset	before	after
8	4	2	end				
9	2	1	CLR	null			
10	9	1	end				

Dirty Page Table

pageID	recLSN

Disk Log

LSN	prvLSN	transID	type	pageID	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				

# Stable Storage – After undo

Disk State

-1 xxxxxxxxxxxx

2 xxxxxxxxxxxx

3 xxxxxxxxxxxx

-1 xxxxxxxxxxxx

Log State							
LSN	prvLSN	transID	type	pageID or undoNext	offset	before	After
2	-1	1	update	2	0	xxx	one
3	-1	2	update	3	0	xxx	two
4	3	2	comm				